

10-605 - HW 7 - Distributed SGD for Matrix Factorization on Spark

Due: Tuesday, April 14, 2015 23:59 EST (Autolab)

Late submission until: Thursday, April 16, 2015 23:59 EST (Autolab)

TAs: Abhinav Maurya, Yipei Wang

Autolab Admin: Yun Ni

Name: _____

Andrew ID: _____

Guidelines for Answers: Please answer to the point, and do not spend time/space giving irrelevant details. You should not require more space than is provided for each question. If you do, please think whether you can make your argument more pithy, an exercise that can often lead to more insight into the problem. Please state any additional assumptions you make while answering the questions. You need to submit a single tar file on autolab, which should include your report. Please make sure you write the report legibly for grading.

Rules for Student Collaboration: The purpose of student collaboration in solving assignments is to facilitate learning, not to circumvent it. Studying the material in groups is strongly encouraged. It is also allowed to seek help from other students in understanding the material needed to solve a particular homework problem, provided no written notes are taken or shared during group discussions, and provided learning is facilitated, not circumvented. The actual solutions must be written and implemented by each student alone, and the student should be ready to reproduce their solution upon request. The presence or absence of any form of help or collaboration, whether given or received, must be explicitly stated and disclosed in full by all involved, on the first page of their assignment. You may ask clarifying questions on Piazza. However, under no circumstances should you reveal any part of the answer publicly on Piazza or any other public website. Any incidents of plagiarism or collaboration without full disclosure will be handled severely in accordance with [CMU's Policy on Academic Integrity](#). The only form of collaboration allowed is with other students who are currently taking the class, and must be disclosed in the Code of Conduct Declaration. In case of any doubts, feel free to consult with one of the instructors.

Rules for External Help: Some of the homework assignments used in this class may have been used in prior versions of this class, or in classes at other institutions. Avoiding the use of heavily tested assignments detracts from the main purpose of these assignments, which is to reinforce the material and stimulate thinking. Because some of these assignments may have been used before, solutions to them may be available online or from other people. It is explicitly forbidden to use any such sources or to consult people who have solved these problems before. You must solve the homework assignments completely on your own. I will mostly rely on your wisdom and honor to follow this rule. However, if a violation is detected, it will be dealt with harshly.

Code of Conduct Declaration

- Did you receive any help whatsoever from anyone in solving this assignment? Yes / No.
- If you answered *yes*, give full details: _____ (e.g. *Jane explained to me what is asked in Question 3.4*)
- Did you give any help whatsoever to anyone in solving this assignment? Yes / No.
- If you answered *yes*, give full details: _____ (e.g. *I pointed Joe to section 2.3 to help him with Question 2*).

Important Note

You are expected to use Python and Spark for this assignment. You cannot use any libraries beyond those already provided in Python. You can use only the built-in constructs of PySpark and are not allowed to use `mllib` or any other Spark library. It could take hours to run and debug your experiments. Start early.

Abhinav Maurya and Yipei Wang are the contact TAs for this assignment. Please post clarification questions to Piazza, and the instructors can be reached at the following email address: `10605-instructors@cs.cmu.edu`.

Brief Introduction to Spark

Spark is a data science software that allows you to write your data processing code in Scala, Python, or Java. The data is loaded as a *Resilient Distributed Database (RDD)* from either the local filesystem or HDFS. RDDs can be converted into other RDDs using transformations such as `map`, `filter`, `reduceByKey`, etc. The evaluation of RDDs is lazy i.e. the required result won't be evaluated until you explicitly invoke an action indicating that you need the result. This allows Spark to optimize the execution of transformations scheduled on RDDs.

Another useful feature of Spark is in-memory processing. You can specify that you want to cache an RDD in memory if you intend to reuse the RDD through multiple iterations of your data processing job. The full set of *transformations* that convert one RDD into another, and *actions* which force the calculation of a result can be found in the [Spark programming guide](#). The programming guide is also a good introduction to Spark. A more detailed RDD API reference with examples can be found [here](#). If you prefer a lecture, you can try the tutorial from Spark Summit 2013 available [here](#).

Spark on Andrew and AWS

Spark is placed in the `bigML` directory. You need to add the location of Spark binaries to your `PATH` variable to be able to run Spark on the Andrew cluster machines. Add the following line to your `.cshrc` file to access Spark on `unix.andrew.cmu.edu` machines:

```
setenv PATH ${PATH}:/afs/cs.cmu.edu/project/bigML/spark-1.3.0-bin-hadoop2.4/bin
```

Instructions on creating an AWS EMR cluster with Spark preinstalled using a bootstrap action can be found [here](#). The important part of the command is:

```
--bootstrap-actions Path=s3://support.elasticmapreduce/spark/install-spark
```

The bootstrap script installs Spark on all EC2 machines of your EMR cluster.

If you are configuring AWS EMR using the web GUI, you can choose *Custom Action* in the *Bootstrap Actions* step, and specify `s3://support.elasticmapreduce/spark/install-spark` as the location of the custom script to be executed.

Big Matrix Factorization using Spark

In this assignment, we will implement [Large-Scale Matrix Factorization with Distributed Stochastic Gradient Descent](#) (DSGD-MF) in Spark. The paper sets forth a solution for matrix factorization using minimization of sum of local losses. The solution involves dividing the matrix into strata for each iteration and performing sequential stochastic gradient descent within each stratum in parallel. The two losses considered are the plain non-zero square loss and the non-zero square loss with L_2 regularization of parameters \mathbf{W} and \mathbf{H} :-

$$L_{ij} = l(\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j}) = (\mathbf{V}_{ij} - \mathbf{W}_{i*}\mathbf{H}_{*j})^2 \quad (1)$$

$$L_{NZSL} = \sum_{(i,j) \in \mathbf{Z}} L_{ij} \quad (2)$$

$$L_2 = L_{NZSL} + \lambda\{\|\mathbf{W}\|_F^2 + \|\mathbf{H}\|_F^2\} \quad (3)$$

Recall that DSGD-MF is a fully distributed algorithm i.e. both the data matrix \mathbf{V} and factor matrices \mathbf{W} and \mathbf{H} can be carefully split and distributed to multiple workers for parallel computation without communication costs between the workers. Hence, it is a good match for implementation in a distributed in-memory data processing system like Spark. We outline the sequential algorithm and describe the steps needed to make it ready for distributed execution.

(a) Stochastic Gradient Descent for Matrix Factorization (SGD-MF)

In SGD-MF (1), we select a single datapoint and update the corresponding row of \mathbf{W} and column of \mathbf{H} in the direction of negative gradient.

The gradients for L_{NZSL} loss are given as follows:-

$$\frac{\partial}{\partial \mathbf{W}_{i*}} l(\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j}) = -2(\mathbf{V}_{ij} - \mathbf{W}_{i*}\mathbf{H}_{*j})\mathbf{H}_{*j} \quad (4)$$

$$\frac{\partial}{\partial \mathbf{H}_{*j}} l(\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j}) = -2(\mathbf{V}_{ij} - \mathbf{W}_{i*}\mathbf{H}_{*j})(\mathbf{W}_{i*})^T \quad (5)$$

The gradients for L_2 loss are given as follows:-

$$\frac{\partial}{\partial \mathbf{W}_{i*}} l(\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j}) = -2(\mathbf{V}_{ij} - \mathbf{W}_{i*}\mathbf{H}_{*j})\mathbf{H}_{*j} + 2\frac{\lambda}{N_{i*}}(\mathbf{W}_{i*})^T \quad (6)$$

$$\frac{\partial}{\partial \mathbf{H}_{*j}} l(\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j}) = -2(\mathbf{V}_{ij} - \mathbf{W}_{i*}\mathbf{H}_{*j})(\mathbf{W}_{i*})^T + 2\frac{\lambda}{N_{*j}}\mathbf{H}_{*j} \quad (7)$$

(b) Stratified Stochastic Gradient Descent for Matrix Factorization (SSGD-MF)

Note that the calculation of gradient of local loss L_{ij} and its use in updating \mathbf{W} and \mathbf{H} is dependent on the entry \mathbf{V}_{ij} , the i^{th} row of \mathbf{W} : \mathbf{W}_{i*} , and the j^{th} column of \mathbf{H} : \mathbf{H}_{*j} .

Algorithm 1 Stochastic Gradient Descent for Matrix Factorization**Require:** Training indices \mathbf{Z} , Training data \mathbf{V} , randomly initialized \mathbf{W}_0 and \mathbf{H}_0

```

while not converged do
  Select a training point  $(i, j) \in \mathbf{Z}$  uniformly at random
   $\mathbf{W}'_{i*} \leftarrow \mathbf{W}_{i*} - \epsilon_n \frac{\partial}{\partial \mathbf{W}_{i*}} l(\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j})$ 
   $\mathbf{H}'_{*j} \leftarrow \mathbf{H}_{*j} - \epsilon_n \frac{\partial}{\partial \mathbf{H}_{*j}} l(\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j})$ 
   $\mathbf{W}_{i*} \leftarrow \mathbf{W}'_{i*}$ 
   $\mathbf{H}_{*j} \leftarrow \mathbf{H}'_{*j}$ 
end while

```

A pair of elements of a matrix given as (i, j) and (i', j') is interchangeable if $i \neq i'$ and $j \neq j'$. See figure (1) for an example. If two such elements are interchangeable, then the stochastic gradient descent updates involving $\{\mathbf{V}_{ij}, \mathbf{W}_{i*}, \mathbf{H}_{*j}\}$ and $\{\mathbf{V}_{i'j'}, \mathbf{W}_{i'*}, \mathbf{H}_{*j'}\}$ do not depend on each other in any way and can be performed in parallel.

A set of elements of a matrix is interchangeable if any pair of elements in the set is interchangeable. The stochastic gradient descent updates involving the local losses of a set of interchangeable elements are also parallel due to the fact that they depend on disjoint parts of the data matrix \mathbf{V} and the factor matrices \mathbf{W} and \mathbf{H} .


We can generalize the notion of interchangeability from elements of a matrix to blocks of the matrix. Consider I and I' are sets of row indices of \mathbf{V} . Similarly, J and J' are sets of column indices of \mathbf{V} . Matrix blocks IJ and $I'J'$ are interchangeable if $I \cap I' = \emptyset$ and $J \cap J' = \emptyset$, where IJ denotes the cartesian product of sets I and J , and \mathbf{V}_{IJ} denotes a matrix block with some abuse and inelegance of notation. Similar to the case of element exchangeability, the stochastic gradient descent updates involving $\{\mathbf{V}_{IJ}, \mathbf{W}_{I*}, \mathbf{H}_{*J}\}$ and $\{\mathbf{V}_{I'J'}, \mathbf{W}_{I'*}, \mathbf{H}_{*J'}\}$ do not depend on each other in any way and can be performed in parallel.

As in the case of element interchangeability, block interchangeability generalizes from a pair of matrix blocks to a set of matrix blocks. A set of matrix blocks is said to be interchangeable if any pair of those matrix blocks is interchangeable. Stochastic gradient descent updates involving interchangeable matrix blocks can be parallelized since they depend on disjoint parts of the data matrix \mathbf{V} and the factor matrices \mathbf{W} and \mathbf{H} .

Given a matrix block \mathbf{V}_{IJ} and the coupled parameter blocks \mathbf{W}_{I*} and \mathbf{H}_{*J} , we can perform sequential stochastic gradient descent for this disjoint part of the matrix and its parameters without having to worry about any parallel updates that might affect \mathbf{W}_{I*} and \mathbf{H}_{*J} . Thus, we have established a concurrent model for stochastic gradient descent updates for matrix factorization. Interchangeable matrix blocks are also called *strata*, and hence this concurrent algorithm of performing SGD is also called Stratified Stochastic Gradient Descent (SSGD).

(c) Distributed Stochastic Gradient Descent for Matrix Factorization (DSGD-MF)

Having established concurrency in SSGD-MF, here we describe how to parallelize SSGD to yield Distributed Stochastic Gradient Descent algorithm for matrix factorization (DSGD-MF). The algorithm is given in (2). As long as the parameter estimates \mathbf{W} and \mathbf{H} have not converged, we choose a set of strata, which creates disjoint blocks of \mathbf{V} , \mathbf{W} and \mathbf{H} that can be handed over to workers for performing sequential stochastic gradient descent. An example of such a stratification is shown in figure (2). Each worker performs SGD updates for the block of parameters it has been handed and returns the updated parameter blocks. For convenience, during iteration i , we will allow the n^{th} worker to perform m_{ni} updates where m_{ni} is the number of non-zero entries in the data matrix

Figure 1: Element Interchangeability in \mathbf{V}

block $V_{I_n J_n}$ handed to that worker during iteration i . The worker will sequentially go through all the non-zero entries in its block $V_{I_n J_n}$ and perform an SGD update for each entry as explained in SGD-MF earlier.

Experiments

We will be using a subsample of the Netflix dataset for the experimental evaluation. This is a very popular recommendation dataset that has ratings of movies by Netflix users. The subsampled dataset is available at http://www.andrew.cmu.edu/user/amaurya/docs/10605/nf_subsample.csv. You will perform your experimental evaluations on the subsampled Netflix dataset. It is in the triples format:

Algorithm 2 Distributed Stochastic Gradient Descent for Matrix Factorization

Require: Training indices \mathbf{Z} , Training data \mathbf{V} , randomly initialized \mathbf{W}_0 and \mathbf{H}_0 , the number of parallel workers B , the number of factors F

while not converged **do**

 Select a stratum $S = \{(I_1, J_1), (I_2, J_2), \dots, (I_B, J_B)\}$ of B blocks

 {parallel for}

for $b \in 1..B$ **do**

 Get block of data matrix $\mathbf{V}_{I_b J_b}$


 Get blocks of parameter matrices $\mathbf{W}_{I_b*}, \mathbf{H}_{*J_b}$

 Perform m_b SGD updates for parameters \mathbf{W}_{I_b*} and \mathbf{H}_{*J_b}

end for

 Collect the updated parameter blocks from all workers and update \mathbf{W} and \mathbf{H}

end while

Figure 2: A stratum in \mathbf{V}

```

<user_1>,<movie_1>,<rating_11>
...
<user_i>,<movie_j>,<rating_ij>
...
<user_M>,<movie_N>,<rating_ij>

```

Here, `user_i` is an integer ID for the i^{th} user, `movie_j` is an integer ID for the j^{th} movie, and `rating_ij` is the rating given by `user_i` to `movie_j`.

For example, the content of the input file to your script will look like:-

```

1,3114,4
1,608,4
1,1246,4
2,1357,5
2,3068,4
2,1537,4
...
6040,562,5
6040,1096,4
6040,1097,4

```

We will explore various choices for the number of iterations I , the number of workers B , the number of factors F , and the inner SGD step size ϵ_n . The SGD stepsize should be set as $\epsilon_n = (\tau_0 + n)^{-\beta}$ where n is the iteration number, so that it decays with iteration number. We will set $\tau_0 = 100$ and vary ϵ_n by varying β . Please, implement the regularized version of DSGD-MF using $\lambda = 0.1$. λ

should be provided as input to your program. For details of all the parameters that your program should expect as inputs, see the later section *Autolab Submission and Evaluation Details*.

Note that you will not have access to the exact iteration number n to calculate ϵ_n while performing SGD updates within a stratum, since strata-specific SGDs are performed in parallel. You need to use an approximate version of the iteration number to calculate ϵ_n . In particular, you should use $n = n' + \sum_{i,b} m_{bi}$ where $\sum_{i,b} m_{bi}$ is the total number of SGD updates made across all strata in all previous iterations and n' is the number of SGD updates made by the current worker on its stratum so far. Thus, ϵ_n is synchronized for all workers at the end of every iteration, but is allowed to be calculated in a decoupled fashion once a worker starts performing SGD updates for the current iteration on its local stratum.

Also, remember to randomly initialize your factor matrices \mathbf{W} and \mathbf{H} ; do not initialize them to zero matrices.

Please perform the following experiments and attach your plots and answers in the report:-

- Set the number of workers $B = 10$, the number of factors $F = 20$, and $\beta = 0.6$. Plot the reconstruction error L_{NZSL} versus the iteration number $i = 1, 2, \dots, 100$. Explain the trend in your plot in the space provided below.

[10 points]

- Set the number of iterations $I = 30$, the number of factors $F = 20$, and $\beta = 0.6$. Plot the runtime of your Spark code R versus number of workers $B = 2, 3, \dots, 10$ in steps of 1. Please ensure your local machine or Spark cluster can support the number of parallel workers you are requesting. Explain the trend in your plot in the space provided below.

[10 points]

- Set the number of iterations $I = 30$, the number of workers $B = 10$, and $\beta = 0.6$. Plot the reconstruction error L_{NZSL} versus number of factors $F = 10, 20, \dots, 100$ in steps of 10. Explain the trend in your plot in the space provided below.

[10 points]

- Set the number of workers $B = 10$, the number of factors $F = 20$, and the number of iterations $I = 30$. Plot the reconstruction error L_{NZSL} versus $\beta = 0.5, 0.6, \dots, 0.9$ in steps of 0.1. Explain the trend in your plot in the space provided below.

[10 points]

Deliverables

In this assignment, you are required to implement DSGD-MF in Python from scratch such that the Python code can be run using PySpark. You need to test the code on the Netflix dataset, provide plots and explanations for your experiments, and provide answers to the theory questions asked in this assignment. You need to submit your implementations via AutoLab where it will be judged on the matrix reconstruction error, time efficiency, and memory efficiency. You should implement the algorithm by yourself instead of using any existing machine learning toolkit. You should upload your code (including all your function files) along with a report named **andrewid-report.pdf**, which should answer the following questions in the space provided:

(1) Is there any advantage to using DSGD for Matrix Factorization instead of Singular Value Decomposition (SVD) which also finds a matrix decomposition that can be used in recommendation systems?

[5 points]

(2) Explain clearly and concisely your method (used in the code you have written) for creating strata at the beginning of every iteration of the DSGD-MF algorithm.

[5 points]

(3) If you were to implement two versions of DSGD-MF using MapReduce and Spark, do you think you will find a relative speedup factor between MapReduce and Spark implementations, keeping

other parameters like the total number of iterations and number of workers fixed? Which implementation do you think will be faster? Why? If your answer depends on any general optimization tricks related to MapReduce or Spark that you know, please state them as well.

[5 points]

(4) Match the Spark RDD transformations to their descriptions. No explanation required.

(1) coalesce	(a) shuffle data randomly into a given number of partitions
(2) repartition	(b) Pin the RDD in memory for repeated use
(3) groupWith	(c) Reduce all values per key to yield a single aggregate value for each key
(4) cache	(d) Merge multiple RDDs based on common keys
(5) foldByKey	(e) Reduce the number of partitions in the RDD possibly without shuffling data

[5 points]

(5) You need to document your code, upload it to a github repository with a `README.md` file at the root of the repository explaining how to run the code, and provide a link to the github repository. We will provide a Google form link on piazza after the assignment deadline where you can submit a link to your Github repository. The code will be checked for readability and documentation. Since this is an important but subjective aspect, we will provide explanation for any deducted marks.

[10 points]

(6) Answer the questions in the collaboration policy on page 1.

Autolab Submission and Evaluation Details

Your DSGD-MF program should be named `dsgd_mf.py` should be able to run without any issues using the following command:

```
spark-submit dsgd_mf.py <num_factors> <num_workers> <num_iterations> \
    <beta_value> <lambda_value> \
    <inputV_filepath> <outputW_filepath> <outputH_filepath>
```

For example,

```
spark-submit dsgd_mf.py 100 10 50 0.8 1.0 test.csv w.csv h.csv
```

Factor matrices obtained from matrix decompositions are often dense and either tall or fat i.e. one of the dimensions is pretty small. Hence, the comma-delimited format is not particularly inefficient for storing factor matrices. Your factor matrices should be written to the files `w.csv` and `h.csv` (provided as commandline arguments) in the normal, non-sparse CSV format. For example, see the following 4x5 matrix.

```
3.0,6,9,12,15
5,10,15.2,20,25
7,14,21.5,28,35
11,22.2,33,44.8,55
```

You can test your code using three test cases provided at [1](#), [2](#), and [3](#). You need not submit your results from the runs on these inputs. But these should give you a fair idea if your algorithm is working correctly.

You need to finally run your code on AWS using [autolab_train.csv](#) as input, and upload the resulting `w.csv` and `h.csv` output files, as well as the Spark logfile from the run. Use 20 as the number of latent factors, and `beta=0.9`. For the AWS cluster, please use 3 machines each of type `m3.xlarge`. To get the best reconstruction error, please cross-validate and choose the best regularization parameter `lambda`, and upload the `w.csv` and `h.csv` output files using the chosen `lambda`. Unlike the experimental evaluations, you are not restricted in the number of iterations and should run your algorithm for as many iterations as is required for convergence of W and H parameters. This will give you the best reconstruction error on the hidden test set we will use for evaluating your submitted `w.csv` and `h.csv` files on autolab.

You should tar the following items into `hw7.tar` and submit to the homework7 assignment via Autolab:

- `andrewid-report.pdf`
- `dsgd_mf.py`, and all other auxiliary files required by `dsgd_mf.py`
- `w.csv` and `h.csv` output files from running your code on [autolab_train.csv](#).
- Spark logfile named `spark_dsgd.log` from running your code on `autolab_train.csv`. The logfile should have all messages printed to terminal when Spark is run on `autolab_train.csv` using the following command:

```
python eval.pyc /tmp/eval_acc.log spark-submit dsgd_mf.py \
    20 3 100 0.9 <optimal_lambda> autolab_train.csv w.csv h.csv
```

- `eval_acc.log` file generated from running the above command

Tar the files directly using `tar -cvf hw7.tar *.py *.csv *.log andrewid-report.pdf`. Do **NOT** put the above files in a folder and then tar the folder. You do not need to upload any saved temporary files. Please make sure your code is working fine before you submit.

You must submit your homework through Autolab via the Homework7: Matrix Factorization link. You have a total of **10 unpenalized submissions** after which each additional submission will be penalized by an additional 2 percent. As previously communicated, you have a total of 2 grace days for this course, which you can use on this or any other assignment. Your performance will be evaluated, and grades from automated evaluation will be provided immediately.

[30 points]

Grading

The total grade of this assignment is **100 points**. You may contact us to use your highest Autolab submission grade after homework is due.

- Tests on your autolab code carry **30 points**. You will be graded based on the memory usage (10 points) and runtime (10 points) for your code, and your final reconstruction error (10 points).
- Experimental results (plots and explanations) on Netflix dataset: **40 points**.
- The theory questions in the assignment will be graded manually: **20 points**.
- Properly documented public repository of the code on Github: **10 points**.

Total: 100
